

Network-Level Insights into Google-Free Android Operating Systems

Bruno Mirčevski¹[0009–0007–1008–6708] and Second Author¹[1111–2222–3333–4444]

Faculty of Computer Science, Białystok University of Technology, Wiejska 45a,
15-351, Białystok, Poland

Abstract. This paper presents a comparative analysis of privacy in Android operating systems through network traffic inspection. Using a controlled man-in-the-middle setup and traffic decryption techniques, six Android distributions were evaluated: stock Google Android, Graphe-neOS, iodéOS and LineageOS (in 3 variants). The study examined out-bound connections, transmitted identifiers, and network traffic patterns to assess the impact of integrated Google services on user privacy. Results indicate that stock Android maintains the most communication with Google servers, transmitting identifiers and metadata even when idle. In contrast, privacy-focused systems minimize unsolicited traffic and offer stronger user control. The findings demonstrate that alternative Android distributions can substantially improve privacy without sacrificing core functionality, highlighting the potential of open-source ecosystems as viable, privacy-respecting alternatives to Google-dependent mobile environments.

Keywords: Android · Privacy · Network traffic.

1 Introduction

Android is the dominant mobile operating system and acts as a primary gateway to digital world for billions of users. It's crucial for core OS to be a reliable, stable and trustworthy platform, that allows users to run apps and services of their choice. While mobile app privacy has been widely studied [2, 11, 9, 15, 8, 14], less attention is typically paid to the operating system layer and to privileged service frameworks that silently enable additional functionality. Standard unprivileged apps are constrained by Android's permission model, although some attempt to circumvent it in creative ways [10, 7]. In contrast, privileged system apps operate with higher access to the device and OS, which makes them harder to understand and control. Most user-facing privacy controls focus on applications and permissions, yet a substantial amount of data exchange can occur below the app layer, through OS services, vendor components, and Google service frameworks. This work therefore compares the network behavior of multiple Android distributions and their preinstalled/privileged system components under controlled, reproducible conditions.

1.1 Android dependency on privacy-threatening services

While Android’s core is the Android Open Source Project (AOSP), most consumer devices ship with additional proprietary components layered on top. They are often implemented as privileged system apps and service frameworks to provide baseline functionality such as push notifications, location assistance, app distribution/updates, device integrity checks and many more. Users and app developers commonly assume their presence and availability by default.

These components exists for a practical reason. Centralized services can significantly improve usability and efficiency, for example by reducing battery drain compared to per-app background connections, and by providing faster, more accurate location through fused network and sensor-based methods. However, this design also tightly couples many functions into a single service stack, so enabling useful capabilities may implicitly enable others that a user would not choose, such as advertising identifiers, telemetry, and behavioral tracking. In practice, the result is a large proprietary bundle that is difficult to audit and decompose into only the features a user actually wants.

On Android, these baseline capabilities are most commonly delivered through Google’s service stack. Because Google’s business model is strongly advertising-driven, many users may be unwilling to entrust this provider with broad access to device and their data. In principle, users should be able to choose which services they rely on and have confidence that these choices are enforced by the platform’s permission and isolation mechanisms. This demand for transparency and controllability motivates alternative approaches such as microG and GrapheneOS’ sandboxed Google Play. These solutions aim to preserve device functionality and app compatibility while reducing privilege and improving user control over Google-related components.

1.2 Privacy-focused Android operating systems

Privacy-focused Android operating systems (also referred to as *custom ROMs*) are distributions derived from AOSP that aim to reduce data exposure at the OS level by limiting preinstalled privileged components, tightening default settings, and providing more transparent controls over permissions and network access. In contrast to many factory builds that bundle large proprietary service stacks, privacy-oriented systems typically try to (i) minimize default background communication, (ii) reduce the use of persistent identifiers, and (iii) give the user stronger enforcement mechanisms to disable or constrain components that are not strictly required for everyday use.

In this work we focus on three representative approaches. GrapheneOS prioritizes a hardened security model and strong isolation. Google components are not included by default, but can be installed in a dedicated sandboxed mode as ordinary apps, without the special privileges that Google Mobile Services (GMS) typically hold on factory systems. This design targets improved containment and user control while preserving high compatibility when the user decides

to install Google services. GrapheneOS also improves low-level security, for example by replacing Android’s default memory allocator with its own hardened implementation to make memory corruption exploitation harder [13].

LineageOS provides a lightweight, broadly supported AOSP-based system with minimal preinstalled software and extensive device compatibility. It can be deployed in multiple configurations: without Google services, with official proprietary Google Play Services and Play Store, or with microG as an open-source reimplementation of key Google APIs. These configurations represent different trade-offs between privacy, compatibility, and reliance on proprietary services.

Finally, iodéOS is a privacy-oriented distribution based on LineageOS that combines a de-Googled baseline with additional privacy features and preconfigured app ecosystem choices, shipping with microG. For completeness, CalyxOS follows a similar privacy-oriented direction (AOSP-based, microG-enabled, and designed for practical daily use). It was considered for the study, but was excluded because it was not receiving updates during the measurement period.

Privacy-focused Android operating systems do not uniformly improve security. Some strengthen the platform through tighter defaults, hardening, and timely patching, while others may weaken it through slower updates or changes that reduce the strength of certain platform guarantees. In particular, microG support often requires signature spoofing, which weakens Android’s signature-based authenticity checks. Systems that enable it typically try to minimize the risk by restricting spoofing to microG. Some distributions do not support verified boot chain or re-locking the bootloader after installation, which weakens protection against physical tampering.

[Summary table...]

1.3 Challenges of network traffic analysis on Android

Network traffic analysis on Android is difficult by design and that is largely a security benefit. Modern protocols increasingly hide data behind encryption (TLS 1.3, QUIC/HTTP3, encrypted DNS). Even when traffic is captured, compression and binary encodings can make payloads hard to interpret without schemas or additional context. As a result, MITM-based decryption is inherently incomplete. Certificate pinning and anti-tamper mechanisms can block interception or break app functionality. Circumvention of these security mechanisms do not reliably yield a full plaintext view. While emulator-based setups enable highly reproducible traffic capture [1], this work uses a physical device to analyze behavior in a real environment, at the cost of reduced control and reproducibility. We accept that some flows will remain opaque, not all network communication will be revealed or understood. Our focus is on differences between studied Android systems and on general patterns.

In this work we captured traffic using a controlled man-in-the-middle gateway around a physical device. A Linux host acted as the Wi-Fi router and transparently redirected the device’s HTTP(S) traffic through an intercepting proxy,

while QUIC over UDP/443 was blocked to force TCP where decryption is feasible. To enable TLS inspection, the proxy CA was installed and promoted to a system-trusted certificate on the device (requiring root access). Where certificate pinning prevented interception, we attempted runtime unpinning via dynamic instrumentation, acknowledging that this only worked for a subset of apps. Full packet captures were recorded and compared across Android distributions to analyze differences in network communication.

2 Related work

Prior work has established that Android privacy risks arise not only from user-installed applications but also, in some cases, from the operating system and its preinstalled components. Privacy-focused Android distributions have gained popularity as alternatives to stock vendor firmware, yet they remain comparatively understudied in the measurement literature. Consequently, there is limited empirical understanding of what practical privacy and control benefits different Android distributions provide to users.

A key OS-level measurement study addresses this gap by comparing several vendor-customized Android variants with two open-source distributions under a privacy-conscious baseline configuration [5]. It reports that all tested stock/vendor builds transmit substantial data to OS vendors and third parties even when the handset is idle, and that this collection persists without an effective opt-out. In contrast, the privacy-focused open-source /e/OS in their comparison exhibits minimal data transmission, whereas LineageOS still shows substantial communication to Google because it was evaluated with Google’s proprietary app stack present. This distinction is important for studies like ours that evaluate multiple distributions and configurations (including variants with and without integrated Google services), since the presence of Google system components can dominate observed traffic and shape the overall privacy profile.

Follow-up work focusing specifically on OEM telemetry further shows that Samsung, Xiaomi, Huawei and Realme make extensive use of long-lived hardware identifiers and also collect installed-app lists and analytics data, sometimes even in connections that should not require persistent identifiers, enabling straightforward linkage to a user’s identity when an OEM account is used [6].

Additional work by the same authors analyzes telemetry from two core Google system apps, Google Messages and Google Dialer, by decrypting and decoding the logging channels that forward data to Google [4]. The study reports that these apps disclose sensitive communication metadata, including when SMS messages and calls are sent or received, with timestamps and for calls duration. Messages also sends a hash derived from message content that can uniquely identify messages and link participants. It further finds that phone numbers may be transmitted and that events are tagged with persistent identifiers such as the Android ID, undermining anonymity and leaving users with no effective opt out. Taken together, these results suggest that even basic default communication

apps on Google Android cannot be assumed to be trustworthy from a privacy perspective.

Complementing OS-level telemetry measurements, a recent study focuses specifically on what pre-installed Google components persistently store on-device and shows that Google's role in Android privacy extends beyond outbound traffic [3]. It reports that Google Play Services, the Play Store, and other bundled Google apps receive and store multiple cookies, advertising-related identifiers, and tracking links even after a factory reset and even when the user has not opened Google apps, with no dedicated consent prompt and no practical opt-out. The study highlights the Google Android ID as a persistent identifier provisioned early and widely reused in subsequent Google connections, and documents how additional tokens and cookies can be stored and later transmitted alongside telemetry or analytics events. Together, these findings reinforce that Google's integrated service stack functions as a central privacy-relevant layer of the Android ecosystem, shaping device identification and data handling independently of user-installed applications.

A broader view of Google's privacy impact is given in a report that describes data collection as an ecosystem across platforms, apps, and advertising services, not as a single feature of Android [12]. It distinguishes "active" data sharing (when users directly use Google products) from "passive" collection that happens in the background via Android/Chrome and via tracking/analytics components embedded in many third-party apps and websites. The report argues that these different data sources can be combined to build detailed user profiles, and that "pseudonymous" identifiers (like ad IDs or cookies) can still be linked back to user accounts when they appear together in the same workflows.

3 Experiment design

4 Results

5 Conclusion

References

1. Jimenez-Berenguel, A., Campo, C., Moure-Garrido, M., Garcia-Rubio, C., Díaz-Sánchez, D., Almenares, F.: Parrot: Portable android reproducible traffic observation tool (Sep 2025). <https://doi.org/10.48550/arXiv.2509.09537>
2. Jin, H., Liu, M., Dodhia, K., Li, Y., Srivastava, G., Fredrikson, M., Agarwal, Y., Hong, J.I.: Why are they collecting my data? inferring the purposes of network traffic in mobile apps. Proc. ACM Interact. Mob. Wearable Ubiquitous Technol. **2**(4) (Dec 2018). <https://doi.org/10.1145/3287051>, <https://doi.org/10.1145/3287051>
3. Leith, D.J.: Cookies, identifiers and other data that google silently stores on android handsets. Tech. rep., Trinity College Dublin (Feb 2025), https://www.scss.tcd.ie/doug.leith/pubs/cookies_identifiers_and_other_data.pdf

4. Leith, D.J.: What data do the google dialer and messages apps on android send to google? In: Li, F., Liang, K., Lin, Z., Katsikas, S.K. (eds.) Security and Privacy in Communication Networks. pp. 549–568. Springer Nature Switzerland, Cham (2023)
5. Liu, H., Patras, P., Leith, D.J.: Android mobile os snooping by samsung, xiaomi, huawei and realme handsets (Oct 2021), https://www.scss.tcd.ie/doug.leith/Android_privacy_report.pdf
6. Liu, H., Patras, P., Leith, D.J.: On the data privacy practices of android oems. PLOS ONE **18**(1), 1–15 (01 2023). <https://doi.org/10.1371/journal.pone.0279942>
7. LocalMess: Disclosure: Covert web-to-app tracking via localhost on android (2025), <https://localmess.github.io/>
8. Pham, A., Dacosta, I., Losiouk, E., Stephan, J., Huguenin, K., Hubaux, J.P.: HideMyApp: Hiding the presence of sensitive apps on android. In: 28th USENIX Security Symposium (USENIX Security 19). pp. 711–728. USENIX Association, Santa Clara, CA (Aug 2019), <https://www.usenix.org/conference/usenixsecurity19/presentation/pham>
9. Razaghpanah, A., Nithyanand, R., Vallina-Rodriguez, N., Sundaresan, S., Allman, M., Kreibich, C., Gill, P.: Apps, trackers, privacy, and regulators: A global study of the mobile tracking ecosystem (01 2018). <https://doi.org/10.14722/ndss.2018.23009>
10. Reardon, J., Feal, Á., Wijesekera, P., On, A.E.B., Vallina-Rodriguez, N., Egelman, S.: 50 ways to leak your data: An exploration of apps' circumvention of the android permissions system. In: 28th USENIX Security Symposium (USENIX Security 19). pp. 603–620. USENIX Association, Santa Clara, CA (Aug 2019), <https://www.usenix.org/conference/usenixsecurity19/presentation/reardon>
11. Ren, J., Lindorfer, M., Dubois, D.J., Rao, A., Choffnes, D.R., Vallina-Rodríguez, N.: Bug fixes, improvements, ... and privacy leaks - a longitudinal study of pii leaks across android app versions. In: Network and Distributed System Security Symposium (2018), <https://api.semanticscholar.org/CorpusID:4231807>
12. Schmidt, D.C.: Google data collection. Tech. rep., Digital Content Next (DCN) (Aug 2018), <https://digitalcontentnext.org/wp-content/uploads/2018/08/DCN-Google-Data-Collection-Paper.pdf>
13. Stefanski, N.: Exploring GrapheneOS secure allocator: Hardened malloc (Sep 2025), <https://www.synacktiv.com/en/publications/exploring-grapheneos-secure-allocator-hardened-malloc>
14. Sutter, T., Tellenbach, B.: Firmwaredroid: Towards automated static analysis of pre-installed android apps (May 2023). <https://doi.org/10.1109/MOBILSoft59058.2023.00009>
15. Wang, Z., Li, Z., Xue, M., Tyson, G.: Exploring the eastern frontier: A first look at mobile app tracking in china. In: Sperotto, A., Dainotti, A., Stiller, B. (eds.) Passive and Active Measurement (PAM 2020). pp. 314–328. Lecture Notes in Computer Science, Springer International Publishing, Cham (2020)